

Team7_project2

November 25, 2025

```
[7]: # Task 1 - Data Preprocessing
import pandas as pd
df = pd.read_csv('proj2_data.csv')
df.head()
```

```
[7]:   InvoiceNo StockCode          Description  Quantity \
0      536365    85123A  WHITE HANGING HEART T-LIGHT HOLDER      6
1      536365      71053           WHITE METAL LANTERN      6
2      536365    84406B    CREAM CUPID HEARTS COAT HANGER      8
3      536365    84029G  KNITTED UNION FLAG HOT WATER BOTTLE      6
4      536365    84029E    RED WOOLLY HOTTIE WHITE HEART.      6

          InvoiceDate  UnitPrice  CustomerID        Country
0  12/1/2010 8:26      2.55    17850.0  United Kingdom
1  12/1/2010 8:26      3.39    17850.0  United Kingdom
2  12/1/2010 8:26      2.75    17850.0  United Kingdom
3  12/1/2010 8:26      3.39    17850.0  United Kingdom
4  12/1/2010 8:26      3.39    17850.0  United Kingdom
```

```
[8]: # Data Preprocessing continued: check missing values
print(df.info())
print(df.describe())
print("The number of missing values is:",
      df.isnull().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   InvoiceNo        541909 non-null   object 
 1   StockCode         541909 non-null   object 
 2   Description       540455 non-null   object 
 3   Quantity          541909 non-null   int64  
 4   InvoiceDate       541909 non-null   object 
 5   UnitPrice         541909 non-null   float64
 6   CustomerID        406829 non-null   float64
 7   Country           541909 non-null   object 
```

```
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
None
      Quantity      UnitPrice   CustomerID
count  541909.000000  541909.000000  406829.000000
mean     9.552250     4.611114  15287.690570
std    218.081158    96.759853  1713.600303
min   -80995.000000 -11062.060000 12346.000000
25%     1.000000     1.250000  13953.000000
50%     3.000000     2.080000  15152.000000
75%    10.000000     4.130000  16791.000000
max    80995.000000  38970.000000 18287.000000
The number of missing values is: InvoiceNo          0
StockCode          0
Description        1454
Quantity           0
InvoiceDate        0
UnitPrice          0
CustomerID        135080
Country            0
dtype: int64
```

```
[9]: # Data Preprocessing continued: replace and check missing values
df['CustomerID'] = df['CustomerID'].fillna('Unknown')
df['Description'] = df['Description'].fillna('Unknown')
print(df.isnull().sum())
```

```
InvoiceNo      0
StockCode      0
Description    0
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    0
Country        0
dtype: int64
```

```
[10]: # Data Preprocessing continued: remove duplicate rows
duplicates = df.duplicated()
total_duplicates = df.duplicated().sum()
print(total_duplicates)
```

```
5268
```

```
[11]: # Data preprocessing: defined the new, cleaned dataframe
df = df.drop_duplicates()
```

```
[12]: # Data Preprocessing: convert InvoiceDate to datetime format
import datetime as dt

df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])

analysis_date = df['InvoiceDate'].max() + pd.Timedelta(days=1)
```

```
[13]: # Guidelines for Analysis - Number 1: Data Overview
print("Dataset Size ")
print("Rows:", df.shape[0])
print("Columns:", df.shape[1])
print("\n Columns Description ")
print(df.dtypes)
print("\n Date Range ")
print("Earliest date:", df['InvoiceDate'].min())
print("Latest date:", df['InvoiceDate'].max())
```

Dataset Size

Rows: 536641

Columns: 8

Columns Description	
InvoiceNo	object
StockCode	object
Description	object
Quantity	int64
InvoiceDate	datetime64[ns]
UnitPrice	float64
CustomerID	object
Country	object
dtype:	object

Date Range

Earliest date: 2010-12-01 08:26:00

Latest date: 2011-12-09 12:50:00

Guidelines for Analysis 1 continued:

1. There are 4373 rows of unique customers and 8 columns.
2. InvoiceNo, StockCode, Description, CustomerID, and Country are all Object data types. Invoice Date was converted to datetime format. Quantity is an integer, and UnitPrice is a float data type. InvoiceNo represents the identification for the purchase, and the stock code represents the identification for what product was purchased. The description is a brief description of the product. Quantiy is the amount of the product that was purchased. InvoiceDate includes the date and time of the purchase. UnitPrice is the price of one unit of the product. CustomerID is the identification for customers and country is the customer's home country.
3. Our dataset covers the period from December 1, 2010 to December 9, 2011.

```
[14]: # Guidelines for Analysis - Number 2: Customer Analysis
print(" Unique Customers ")
unique_customers = df['CustomerID'].nunique()
print(unique_customers)
print("\n Orders per Customer ")
orders_per_customer = df.groupby('CustomerID')['InvoiceNo'].nunique()
print(orders_per_customer.describe())
print("\n Top 5 Customers by Order Count ")
print(orders_per_customer.sort_values(ascending=False).head())
```

Unique Customers

4373

Orders per Customer

count	4373.000000
mean	5.922708
std	56.798813
min	1.000000
25%	1.000000
50%	3.000000
75%	5.000000
max	3710.000000

Name: InvoiceNo, dtype: float64

Top 5 Customers by Order Count

CustomerID	
Unknown	3710
14911.0	248
12748.0	224
17841.0	169
14606.0	128

Name: InvoiceNo, dtype: int64

```
[15]: # Guidelines for Analysis - Number 3: Product Analysis
```

```
print(" Top 10 Most Purchased Products ")
top_products = df.groupby('Description')['Quantity'].sum() .
    sort_values(ascending=False)
print(top_products)
print("\n Average Unit Price of Products ")
print(df['UnitPrice'].mean())
print("\n Revenue By Product ")
df['Revenue'] = df['Quantity'] * df['UnitPrice']
top_revenue = df.groupby('Description')['Revenue'].sum() .
    sort_values(ascending=False)
print(top_revenue)
```

Top 10 Most Purchased Products

Description

WORLD WAR 2 GLIDERS ASSTD DESIGNS	53751
JUMBO BAG RED RETROSPOT	47260
POPCORN HOLDER	36322
ASSORTED COLOUR BIRD ORNAMENT	36282
PACK OF 72 RETROSPOT CAKE CASES	36016
	...
Printing smudges/thrown away	-9058
check	-12030
Unknown	-13609
Unsaleable, destroyed.	-15644
printing smudges/thrown away	-19200

Name: Quantity, Length: 4221, dtype: int64

Average Unit Price of Products
4.632655674836623

Revenue By Product

Description

DOTCOM POSTAGE	206245.480
REGENCY CAKESTAND 3 TIER	164459.490
WHITE HANGING HEART T-LIGHT HOLDER	99612.420
PARTY BUNTING	98243.880
JUMBO BAG RED RETROSPOT	92175.790
	...
Bank Charges	-7175.639
CRUK Commission	-7933.430
Adjust bad debt	-11062.060
Manual	-69031.640
AMAZON FEE	-221520.500

Name: Revenue, Length: 4221, dtype: float64

[16]: # Guidelines for Analysis - Number 4: Time Analysis

```
df['DayOfWeek'] = df['InvoiceDate'].dt.day_name()
df['Hour'] = df['InvoiceDate'].dt.hour
df['Month'] = df['InvoiceDate'].dt.to_period('M')
print(" Orders by Day of Week ")
print(df['DayOfWeek'].value_counts())
print("\n Orders by Hour ")
print(df['Hour'].value_counts().sort_index())
print("\n Monthly Orders (Seasonality) ")
print(df['Month'].value_counts().sort_index())
```

Orders by Day of Week

DayOfWeek	
Thursday	103056
Tuesday	101064

```
Monday      94435
Wednesday   93715
Friday      81565
Sunday      62806
Name: count, dtype: int64
```

```
Orders by Hour
Hour
6          41
7         383
8        8906
9       34314
10      48808
11      56949
12      77573
13      71247
14      66572
15      76938
16      54134
17      28371
18      7941
19      3617
20      847
Name: count, dtype: int64
```

```
Monthly Orders (Seasonality)
Month
2010-12    41981
2011-01    34900
2011-02    27479
2011-03    36439
2011-04    29701
2011-05    36782
2011-06    36609
2011-07    39267
2011-08    35064
2011-09    49861
2011-10    59969
2011-11    83343
2011-12    25246
Freq: M, Name: count, dtype: int64
```

```
[17]: # Guidelines for Analysis - Number 5: Geographical Analysis
import matplotlib.pyplot as plt
import seaborn as sns

print(" Top 5 Countries by Order Count ")
```

```

top_countries = df.groupby('Country')['InvoiceNo'].nunique() .
    ↪sort_values(ascending=False)
print(top_countries.head())
print("\n Average Order Value by Country ")
country_aov = df.groupby('Country')['Revenue'].mean() .
    ↪sort_values(ascending=False)
print(country_aov.head())

import matplotlib.pyplot as plt
import seaborn as sns

country_orders = df.groupby('Country')['InvoiceNo'].nunique()

country_aov = df.groupby('Country')['Revenue'].mean()

country_stats = (
    pd.concat([country_orders, country_aov], axis=1)
    .reset_index()
)
country_stats.columns = ['Country', 'NumOrders', 'AvgOrderValue']

country_stats = country_stats[country_stats['Country'] != 'United Kingdom']

# Plot
plt.figure(figsize=(12, 7))
sns.scatterplot(
    data=country_stats,
    x='NumOrders',
    y='AvgOrderValue',
    s=100
)

for _, row in country_stats.iterrows():
    plt.text(row['NumOrders'] + 0.1, row['AvgOrderValue'], row['Country'], ↪
    fontsize=8)

plt.title('Number of Orders vs. Average Order Value by Country', fontsize=16)
plt.xlabel('Number of Orders')
plt.ylabel('Average Order Value')
plt.tight_layout()
plt.show()

```

Top 5 Countries by Order Count

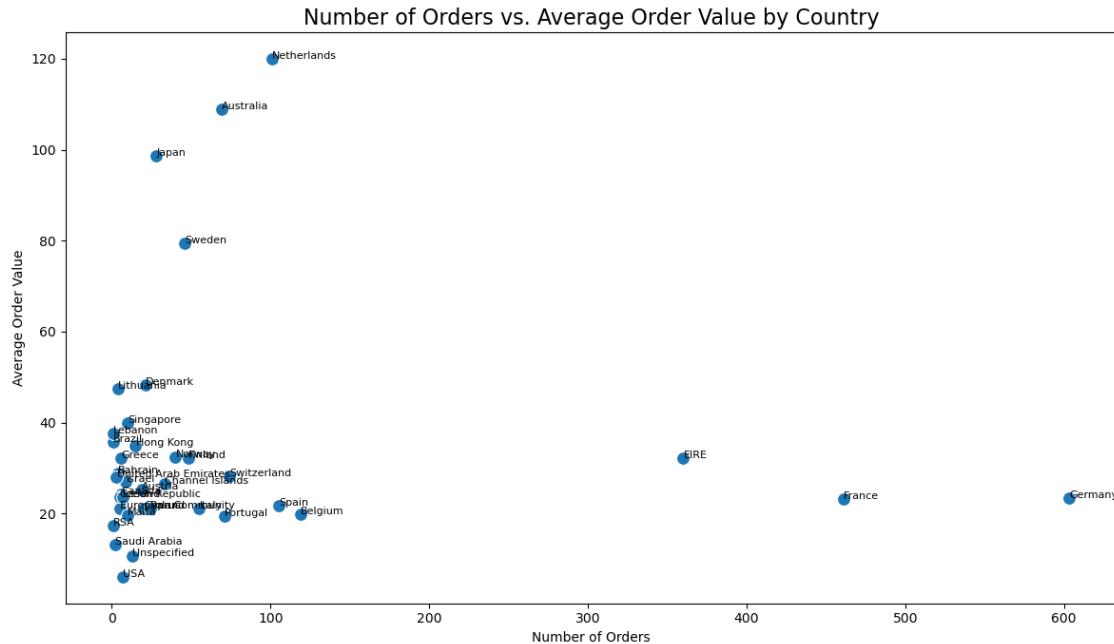
Country	
United Kingdom	23494
Germany	603
France	461

```
EIRE          360
Belgium      119
Name: InvoiceNo, dtype: int64
```

Average Order Value by Country

Country	Average Order Value
Netherlands	120.059696
Australia	108.910787
Japan	98.716816
Sweden	79.360976
Denmark	48.247147

```
Name: Revenue, dtype: float64
```



Geographical Analysis continued:

Outside of the United Kingdom, there appears to be a slight negative correlation between average order value and the number of orders per country.

```
[18]: # Geographical Analysis continued
import matplotlib.pyplot as plt
import seaborn as sns

# Create a DataFrame from the Series for easier plotting
country_aov_df = country_aov.reset_index()
country_aov_df.columns = ['Country', 'AverageOrderValue']

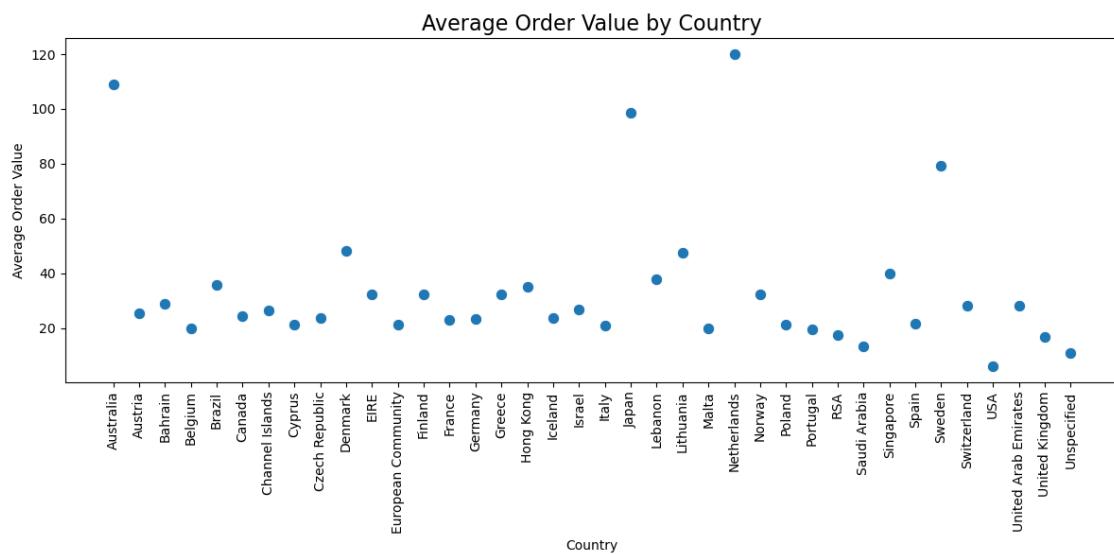
plt.figure(figsize=(12,6))
```

```

sns.scatterplot(
    data=country_aov_df,
    x='Country',
    y='AverageOrderValue',
    s=80
)

plt.title('Average Order Value by Country', fontsize=16)
plt.xlabel('Country')
plt.ylabel('Average Order Value')
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()

```



Guidelines for Analysis: Geographical Analysis continued:

Other than the possible slight negative correlation between average order value and the number of orders per country, there seems to be no correlation between customer's country and average order value.

[19]: # Guidelines for Analysis - Number 6: Payment Analysis

```

print("Dataset does NOT contain payment method - cannot analyze payment method trends.")

```

Dataset does NOT contain payment method - cannot analyze payment method trends.

[20]: # Task 2 - RFM Calculation

```

rfm = df.groupby('CustomerID').agg({
    'InvoiceDate': lambda x: (analysis_date - x.max()).days,
})

```

```

    'InvoiceNo': 'nunique',
    'UnitPrice': 'sum'
})

rfm.rename(columns={
    'InvoiceDate': 'Recency',
    'InvoiceNo': 'Frequency',
    'UnitPrice': 'Monetary'
}, inplace=True)

rfm.head(10)

```

[20]:

	CustomerID	Recency	Frequency	Monetary
12346.0	326	2	2.08	
12347.0	2	7	481.21	
12348.0	75	4	178.71	
12349.0	19	1	605.10	
12350.0	310	1	65.30	
12352.0	36	11	2211.10	
12353.0	204	1	24.30	
12354.0	232	1	261.22	
12355.0	214	1	54.65	
12356.0	23	3	188.87	

[21]:

```

# Task 3 - RFM Segmentation
rfm['Recency_rank'] = rfm['Recency'].rank(method='first')
rfm['Frequency_rank'] = rfm['Frequency'].rank(method='first')
rfm['Monetary_rank'] = rfm['Monetary'].rank(method='first')

rfm['R_score'] = pd.qcut(rfm['Recency_rank'], 4, labels=[4,3,2,1]).astype(int)
rfm['F_score'] = pd.qcut(rfm['Frequency_rank'], 4, labels=[1,2,3,4]).astype(int)
rfm['M_score'] = pd.qcut(rfm['Monetary_rank'], 4, labels=[1,2,3,4]).astype(int)

rfm['RFM'] = (
    rfm['R_score'] +
    rfm['F_score'] +
    rfm['M_score']
)
print(rfm['RFM'])

```

CustomerID	RFM
12346.0	4
12347.0	12
12348.0	8
12349.0	8
12350.0	4
	..

```
18281.0      4
18282.0      9
18283.0     12
18287.0      8
Unknown       12
Name: RFM, Length: 4373, dtype: int64
```

```
[22]: # Also included are the individual scores for each category
rfm_features = rfm[['R_score', 'F_score', 'M_score']]
rfm_features
```

```
[22]:      R_score  F_score  M_score
CustomerID
12346.0          1         2         1
12347.0          4         4         4
12348.0          2         3         3
12349.0          3         1         4
12350.0          1         1         2
...
...           ...   ...   ...
18281.0          1         2         1
18282.0          4         3         2
18283.0          4         4         4
18287.0          3         3         2
Unknown          4         4         4
```

[4373 rows x 3 columns]

```
[26]: !pip install kneed
```

```
Collecting kneed
  Using cached kneed-0.8.5-py3-none-any.whl.metadata (5.5 kB)
Requirement already satisfied: numpy>=1.14.2 in
/opt/anaconda3/lib/python3.12/site-packages (from kneed) (2.1.3)
Requirement already satisfied: scipy>=1.0.0 in
/opt/anaconda3/lib/python3.12/site-packages (from kneed) (1.16.3)
Using cached kneed-0.8.5-py3-none-any.whl (10 kB)
Installing collected packages: kneed
Successfully installed kneed-0.8.5
```

```
[27]: # Task 4 - Customer Segmentation
# Use K-Means to find suggested optimal number of clusters
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
from kneed import KneeLocator

scaler = StandardScaler()
rfm_scaled = scaler.fit_transform(rfm_features)
```

```

wcss = []

for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(rfm_scaled)
    wcss.append(kmeans.inertia_)

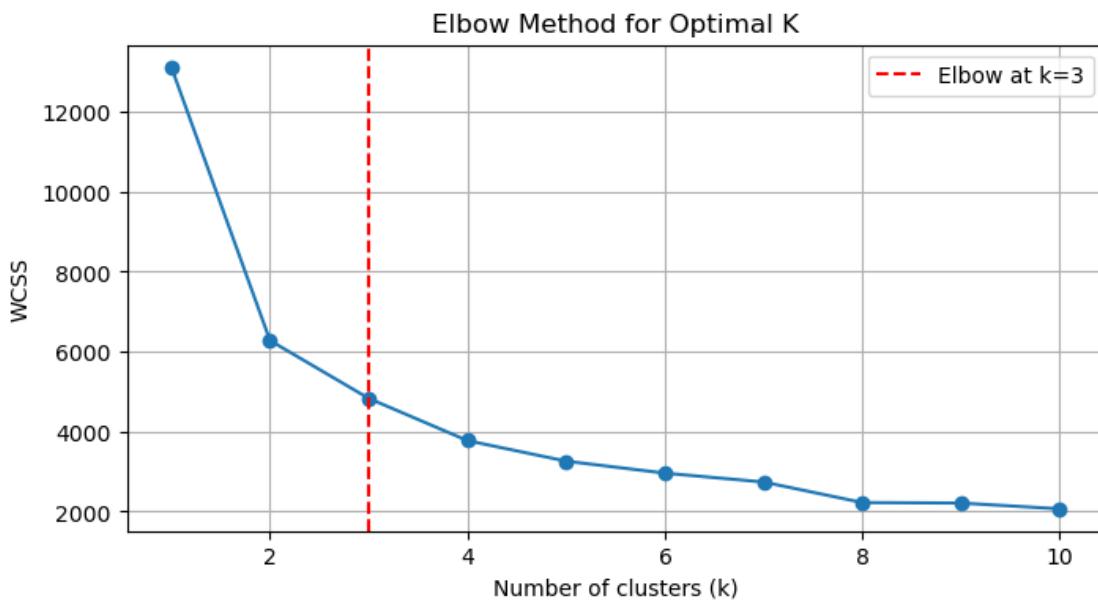
plt.figure(figsize=(8,4))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method for Optimal K')
plt.xlabel('Number of clusters (k)')
plt.ylabel('WCSS')
plt.grid(True)

kl = KneeLocator(range(1, 11), wcss, curve='convex', direction='decreasing')
optimal_k = kl.elbow
plt.axvline(x=optimal_k, color='r', linestyle='--', label=f'Elbow at k={optimal_k}')
plt.legend()

plt.show()

print(f"The estimated optimal number of clusters is: {optimal_k}")

```



The estimated optimal number of clusters is: 3

```
[28]: # Customer Segmentation continued: experiment with 2 clusters
from sklearn.cluster import KMeans

k = 2

kmeans = KMeans(n_clusters=k, random_state=11)
rfm['Cluster'] = kmeans.fit_predict(rfm_scaled)

cluster_summary = rfm.groupby('Cluster')[['R_score', 'F_score', 'M_score']].mean()
print(cluster_summary)
```

Cluster	R_score	F_score	M_score
0	3.208498	3.453557	3.426383
1	1.890166	1.677735	1.701149

```
[29]: # Customer Segmentation continued: experiment with 3 clusters
k = 3

kmeans = KMeans(n_clusters=k, random_state=12)
rfm['Cluster'] = kmeans.fit_predict(rfm_scaled)

cluster_summary = rfm.groupby('Cluster')[['R_score', 'F_score', 'M_score']].mean()
print(cluster_summary)
```

Cluster	R_score	F_score	M_score
0	3.225187	3.529883	3.500534
1	1.000000	1.625731	1.750487
2	2.623218	1.797692	1.748133

```
[30]: # Customer Segmentation continued: experiment with 4 clusters
k = 4

kmeans = KMeans(n_clusters=k, random_state=13)
rfm['Cluster'] = kmeans.fit_predict(rfm_scaled)

cluster_summary = rfm.groupby('Cluster')[['R_score', 'F_score', 'M_score']].mean()
print(cluster_summary)
```

Cluster	R_score	F_score	M_score
0	1.404454	1.465517	1.503592
1	1.900208	2.980249	3.047817
2	3.359712	1.932854	1.821343
3	3.670042	3.723207	3.702110

```
[31]: # Customer Segmentation continued: 3 clusters selected as optimal number of clusters
k = 3

kmeans = KMeans(n_clusters=k, random_state=12)
rfm['Cluster'] = kmeans.fit_predict(rfm_scaled)

cluster_summary = rfm.groupby('Cluster')[['R_score', 'F_score', 'M_score']].mean()
print(cluster_summary)

cluster_counts = rfm.groupby('Cluster').size()
print(cluster_counts)
```

Cluster	R_score	F_score	M_score
0	3.225187	3.529883	3.500534
1	1.000000	1.625731	1.750487
2	2.623218	1.797692	1.748133

Cluster	
0	1874
1	1026
2	1473

dtype: int64

5. Segment Profiling

- We will use data from $k = 3$ clusters, as this data was the most clear. It also supports our findings from the elbow method for optimal k .
- Cluster 0 has the highest means of all three clusters. This means that Cluster 0 contains the most recent customers, the most frequent customers, and those customers that spent the most money. This cluster contains the fewest customers.
- On the other end, Cluster 2 represents the customers that have not purchased recently, are not frequent buyers, and who spent little money on their purchases. This cluster contains the most customers.
- Cluster 1 represents the middle cluster, showing customers who spend an average amount of money and are neither frequent nor infrequent customers, and purchase at an average recency.

6. Marketing Recommendations

Cluster 0: For our most frequent, recent, and highest spending customers, we should create a loyalty or VIP program. This could include early access to products, especially the nicer, more expensive, or more exclusive products (or make some products VIP exclusive). VIP Members could also be sent small letters of thanks or small gifts, particularly around the holidays or the customers' birthday. Maintaining these customers is important because it helps maintain a steady source of revenue.

Cluster 1: For our ‘average’ customers, we should view them as potential candidates to become VIPs. To do so, we could create limited-time offers on more exclusive products and offer incentives such as ‘Spend an additional \$25 to get a free gift’. We can also send them

notifications on how close they are to VIP status and the benefits of reaching that status. This strategy seeks to retain customers and persuade them to reach VIP status, which would generate higher revenue.

Cluster 2: For our rare customers, we should strive to generate more attention to our company. We can focus on marketing strategies such as free shipping for a limited time only or discounts on some select items. This should generate more interest in our company and could attract new customers as well.

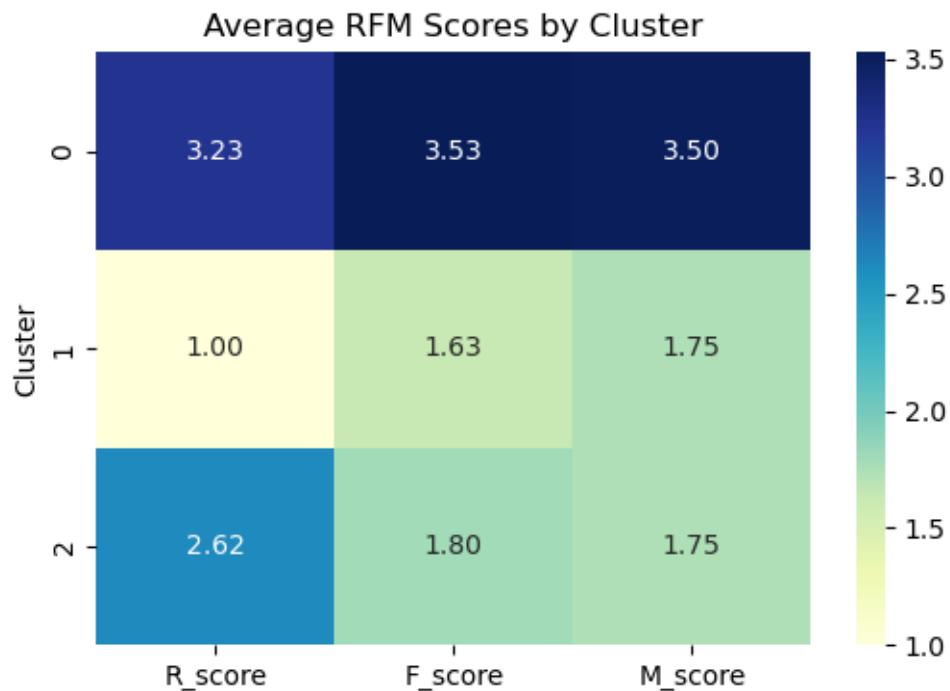
```
[32]: # Task 7 - Visualizations  
import seaborn as sns
```

```
plt.figure(figsize=(8,6))  
sns.scatterplot(x='Recency_rank', y='Frequency_rank', hue='Cluster', data=rdf,  
                 palette='Set1', s=60)  
plt.title('Recency vs Frequency by Cluster')  
plt.xlabel('Recency Rank')  
plt.ylabel('Frequency Rank')  
plt.legend(title='Cluster')  
plt.show()
```

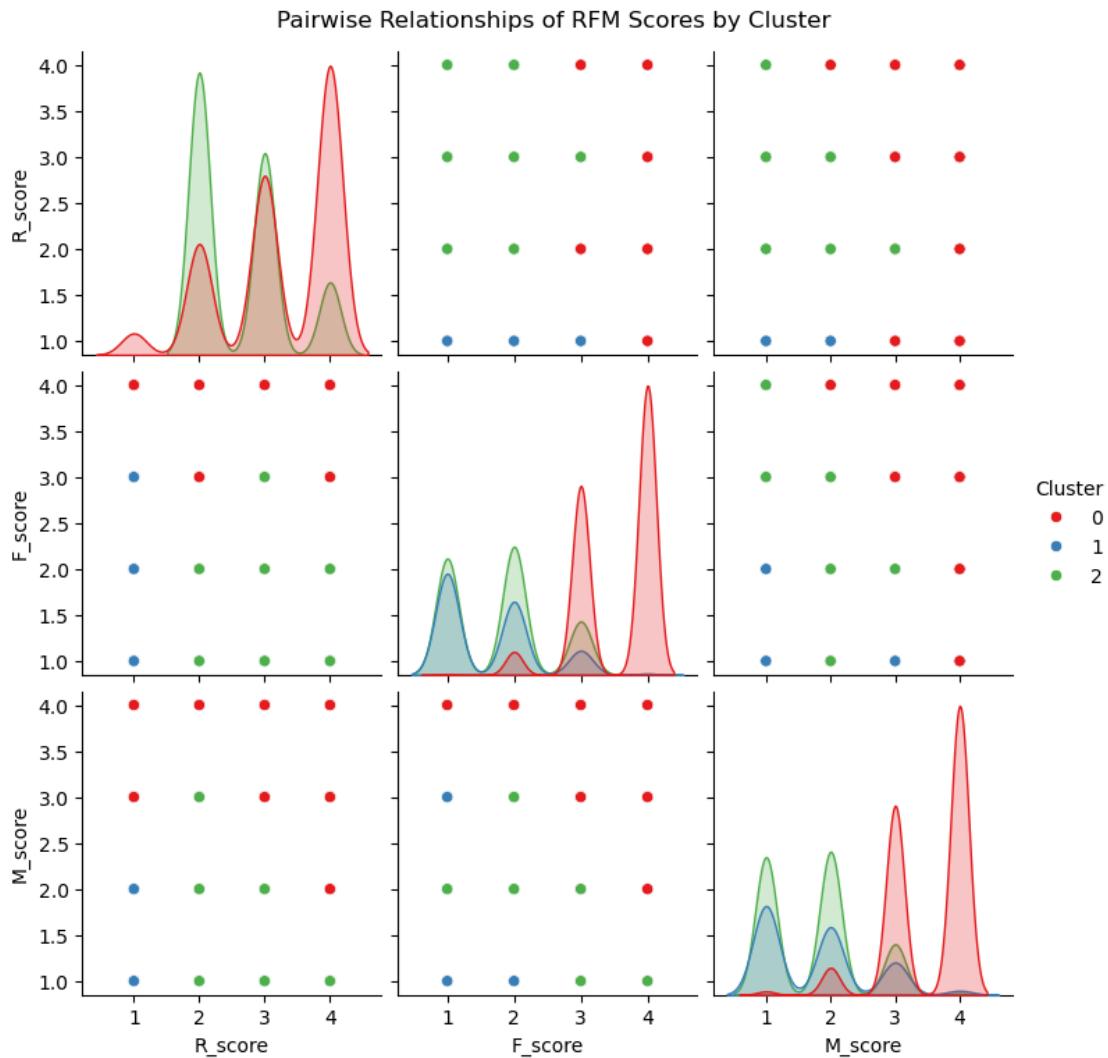


```
[33]: # Visualizations continued
cluster_summary = rfm.groupby('Cluster')[['R_score','F_score','M_score']].mean()

plt.figure(figsize=(6,4))
sns.heatmap(cluster_summary, annot=True, cmap='YlGnBu', fmt=".2f")
plt.title('Average RFM Scores by Cluster')
plt.show()
```



```
[34]: # Visualizations continued
sns.pairplot(rfm, vars=['R_score','F_score','M_score'], hue='Cluster', palette='Set1', diag_kind='kde')
plt.suptitle('Pairwise Relationships of RFM Scores by Cluster', y=1.02)
plt.show()
```



[35]: # Guidelines for Analysis - Number 7: Customer Behavior

```

print(" Customer Lifespan (First to Last Purchase) ")
lifespan = df.groupby('CustomerID')[['InvoiceDate']].agg(['min', 'max'])
lifespan['ActiveDays'] = (lifespan['max'] - lifespan['min']).dt.days
print(lifespan['ActiveDays'].describe())

print("\n Sample Lifespan for 10 Customers ")
print(lifespan.head(10))

```

```

Customer Lifespan (First to Last Purchase)
count    4373.000000
mean     133.440430
std      132.879274
min      0.000000

```

```

25%      0.000000
50%      98.000000
75%      256.000000
max      373.000000
Name: ActiveDays, dtype: float64

```

Sample Lifespan for 10 Customers

		min	max	ActiveDays
CustomerID				
12346.0	2011-01-18 10:01:00	2011-01-18 10:17:00		0
12347.0	2010-12-07 14:57:00	2011-12-07 15:52:00		365
12348.0	2010-12-16 19:09:00	2011-09-25 13:13:00		282
12349.0	2011-11-21 09:51:00	2011-11-21 09:51:00		0
12350.0	2011-02-02 16:01:00	2011-02-02 16:01:00		0
12352.0	2011-02-16 12:33:00	2011-11-03 14:37:00		260
12353.0	2011-05-19 17:47:00	2011-05-19 17:47:00		0
12354.0	2011-04-21 13:11:00	2011-04-21 13:11:00		0
12355.0	2011-05-09 13:49:00	2011-05-09 13:49:00		0
12356.0	2011-01-18 09:50:00	2011-11-17 08:40:00		302

Customer Behavior continued:

As mentioned earlier, it appears that each customer segment (cluster) is divided by purchase behavior recency. However, the range of customer activity does overlap across the clusters.

[36]: # Guidelines for Analysis - Number 8: Returns and Refunds

```

returns = df[df['Quantity'] < 0]

print(" Percent of Orders with Returns ")
return_rate = (returns['InvoiceNo'].nunique() / df['InvoiceNo'].nunique()) * 100
print(f"{return_rate:.2f}% of orders include returns.")

print("\n Product Categories with Most Returns ")
returns_by_product = returns.groupby('Description')['Quantity'].sum().
    sort_values().head(10)
print(returns_by_product)

```

Percent of Orders with Returns
19.97% of orders include returns.

Description	Quantity
PAPER CRAFT , LITTLE BIRDIE	-80995
MEDIUM CERAMIC TOP STORAGE JAR	-74494
Unknown	-46156
printing smudges/thrown away	-19200
Unsaleable, destroyed.	-15644
check	-13247

```
? -9496
ROTATING SILVER ANGELS T-LIGHT HLDR -9376
Printing smudges/thrown away -9058
Damaged -7540
Name: Quantity, dtype: int64
```

[37]: # Guidelines for Analysis - Number 9: Profitability

```
print("Dataset does NOT include product cost → cannot compute profit or margin. We can compute Revenue instead.")

# Compute Revenue instead

df['Revenue'] = df['Quantity'] * df['UnitPrice']
total_revenue = total_revenue = round(df['Revenue'].sum(), 2)
print("Total Revenue is: ", total_revenue)
```

Dataset does NOT include product cost → cannot compute profit or margin. We can compute Revenue instead.

Total Revenue is: 9726006.95

[39]: # Guidelines for Analysis - Number 10: Customer Satisfaction

```
print("Dataset includes NO ratings/reviews → customer satisfaction analysis not possible.")
```

Dataset includes NO ratings/reviews → customer satisfaction analysis not possible.

[]: